

Large Survey Database, Continued

But First

- Homework #1 assignment is in the homeworks/hw1 subdirectory of ast597b repository on GitHub!

<http://nbviewer.ipython.org/github/mjuric/ast597b/tree/master/homeworks/hw1/>

*Large **S**urvey **D**atabase*

- A distributed Python framework for storing, querying, and distributing computation on large datasets
- <http://lsddb.org>
- Data access: SQL-like queries (Python)
- Current release:
 - SQL-like query language
 - Local caching of data
 - On-the-fly cross-matching of catalogs
 - ACID (Atomicity, Consistency, Isolation, Durability) transactions
 - MapReduce engine and Python API
 - Multi-core aware query engine
 - Distributed query engine (experimental)

Limitations

- A made-up, non-standard, language
- Limited ability to do joins
- May have strange, buggy, corner cases
- Only a few active developers


LSD @ ASTR 597b

- We will primarily be using LSD to extract subsets from large data sets (mostly SDSS, WISE)
- We will use it to do catalog cross-matching
- We will be importing catalogs to Survey Science Group's LSD server
 - This will be your final project: writing the table schemas and code to load a survey into LSD
 - Note: still procuring the machine, should be here by the end of the quarter.
 - Result: one of the more capable survey archives!

Architecture & Performance

- **LSD is designed to work with simple catalogs; one catalog is one table.**
 - Catalogs with multiple observations are an exception and usually have three tables (one for objects, one for observations, and a “join” table connecting the two)
 - Not optimal, but generally works
- **Each table is spatially (and, possibly temporally) partitioned**
 - The sky is subdivided in cells in space and time
 - Each cell is stored in its own directory, with multiple files containing the data (see below)
 - This is known as *horizontal partitioning* or *sharding* in database terminology
- **Each table is partitioned into a number of “column groups” – groups of columns that are frequently used together**
 - Each column group is stored in its own file within a cell
 - This is known as *vertical partitioning*
- **This design results in better performance than with traditional databases:**
 - Only the data that’s needed for the query will be loaded from disk (i.e., infrequently used column groups won’t be read unless referred to)
 - A query can be subdivided into pieces that run in parallel on multiple cells
 - The parallelization can span multiple machines


Healpix Projected, Spatially Partitioned Data



Pan-STARRS Cluster Network last hour

Example: Sweep through the DB on Odyssey (10,000+ core cluster)

- Almost linear scaling
(from 20 to 36 nodes)
- I/O limited (1Gbps/node;
20Gbps total, switches in
racks also a bottleneck)
- Should scale
to hundreds of nodes



Topics

1. Installing LSD
2. Creating tables and importing data
3. Connecting to remote tables
4. Querying from the command line
5. Catalog cross-matching with LSD
6. Using LSD from Python

In class git repo @ [lectures/2015-02-02-lsd/lecture5.ipynb](#)

Heavy parallelism

```
top - 12:25:23 up 3 days,  2:55, 11 users,  load average: 6.91, 7.38, 5.04
Tasks: 603 total,  18 running, 585 sleeping,   0 stopped,   0 zombie
Cpu(s): 49.0%us,  1.5%sy,  3.0%ni, 46.5%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Mem: 264508168k total, 224618492k used, 39889676k free,   801192k buffers
Swap: 32767996k total,    8472k used, 32759524k free, 185778252k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
20258	mjuric	20	0	2086m	405m	4328	R	107.9	0.2	0:23.25	lsd-import
20268	mjuric	20	0	2185m	1.8g	1908	R	100.0	0.7	0:23.01	lsd-import
20254	mjuric	20	0	3264m	1.8g	4224	R	99.9	0.7	0:23.01	lsd-import
20255	mjuric	20	0	2487m	860m	4260	R	99.9	0.3	0:23.30	lsd-import
20257	mjuric	20	0	3335m	1.7g	4356	R	99.9	0.7	0:22.84	lsd-import
20259	mjuric	20	0	2173m	1.8g	1908	R	99.9	0.7	0:23.01	lsd-import
20261	mjuric	20	0	1932m	1.5g	1932	R	99.9	0.6	0:23.01	lsd-import
20262	mjuric	20	0	3514m	1.8g	4364	R	99.9	0.7	0:23.06	lsd-import
20263	mjuric	20	0	2167m	1.8g	1908	R	99.9	0.7	0:23.00	lsd-import
20266	mjuric	20	0	1905m	1.5g	1908	R	99.9	0.6	0:23.00	lsd-import
20256	mjuric	20	0	1901m	1.5g	1932	R	99.6	0.6	0:23.00	lsd-import
20260	mjuric	20	0	1918m	1.5g	1932	R	99.6	0.6	0:23.00	lsd-import
20264	mjuric	20	0	2156m	1.7g	1908	R	99.6	0.7	0:22.95	lsd-import
20265	mjuric	20	0	1955m	1.5g	1932	R	99.6	0.6	0:22.96	lsd-import
20267	mjuric	20	0	2010m	1.6g	1908	R	99.6	0.6	0:23.00	lsd-import
20269	mjuric	20	0	2349m	1.9g	1908	R	99.6	0.8	0:23.00	lsd-import

LSD Queries

LSD Query Syntax

1.) SQL-like, case-insensitive keywords

2.) Column specifications are Python expressions; free to call Python functions from within query clauses

```
SeLEct
    ra, dec, g, r, sdss.g as sg, sdss.r as sr,
    sg-sr as sgr, ffitkw(chip_hdr(chip_id), "ZPT_OBS") as zpt
FROM
    ps1_obj, ps1_det, ps1_exp, sdss(outer)
WHERE
    sr < 22.5
INTO
    mytable
```

4.) The WHERE clause is a Python expression, with column data given in NumPy arrays.

3.) Implicit natural JOINS between tables, outer JOINS supported